

USING R STAN

R User Group meeting

October 27, 2022

Fayette Klaassen, PhD

TODAY'S TALK

- Brief introduction Bayesian statistics
- Introduction to stan and rstan
- Workflow & troubleshooting rstan
- Demonstration
- Questions

WHAT IS BAYESIAN STATISTICS?

- Bayes theorem

$$\frac{P(D | \theta)P(\theta)}{P(D)} = P(\theta | D)$$

- $P(D | \theta)$ likelihood of the data
 - $P(\theta)$ prior distribution
 - $P(D)$ marginal distribution
 - $P(\theta | D)$ posterior distribution
- Combining prior knowledge about θ to learn the conditional probability of θ given D

WHY BAYESIAN STATISTICS?


- Conditional probability (parameters given data)
 - Instead of frequentist probability of the data given a parameter
- Include prior information
 - Accumulate evidence across research
- Facilitate fitting more complex models (decrease the complexity of a model)

EXAMPLE


- Flipping a coin 10 times
 - 7 heads
- 4 prior distributions
 - Fair coin
 - Very certain fair coin
 - Unknown coin
 - Unfair coin
- Data and prior combined

To posterior


Fair coin prior Beta(10,10)




Very strong prior Beta(100,100)



Uniform prior Beta(1,1)



Unfair coin prior Beta(10,1)



BAYESIAN ESTIMATION

- Posterior distribution often not closed form solution
- Iterative sampling procedure
 - MCMC
- Determine the posterior probability for a set of prior values
 - Choosing priors + initial values wisely!

h	θ_1	θ_2
0	-	$\theta_2^{(0)}$
1	$\theta_1^{(1)} \theta_2^{(0)}$	$\theta_2^{(1)} \theta_1^{(1)}$
2	$\theta_1^{(2)} \theta_2^{(1)}$	$\theta_2^{(2)} \theta_1^{(2)}$
3	$\theta_1^{(3)} \theta_2^{(2)}$	$\theta_2^{(3)} \theta_1^{(3)}$
...		
H	$\theta_1^{(H)} \theta_2^{(H-1)}$	$\theta_2^{(H)} \theta_1^{(H)}$

METHODS FOR BAYESIAN ESTIMATION

Direct specification of posterior	only for simple models
Gibbs sampling (MCMC algorithm)	implemented in JAGS ; BUGS ; requires a closed form of the (conditional) posterior
Metropolis Hastings algorithm	time intensive; explore posterior using proposal distribution
Hamiltonian Monte Carlo	uses gradient of un-normalized log-posterior; more efficient and complete sampling
No U-Turn Sampler	auto-tuning of leapfrog steps
Optimization routines (e.g. BFGS in stan)	maximizing the objective function; no uncertainty intervals

WHAT IS RSTAN

- Stan
 - C++ language for Bayesian modeling
 - Utilizes Hamiltonian Monte Carlo (HMC) and No U-Turn Sampler (NUTS) routines
 - Hoffman & Gelman (2014), and [stan core development team](#)
- Rstan
 - R interface (package) for stan
 - Requires installation of a C++ compiler (gcc/gnu)
 - Requires several R packages
 - Installation can be tricky – follow instructions
 - <https://cran.r-project.org/web/packages/rstan/vignettes/rstan.html>



WHY USE RSTAN

- Get the speed and efficiency from the HMC and NUTS in stan
 - Auto-differentiation for HMC, auto-tuning for NUTS
- Ability to specify and fit complex models
- Get the ease of writing the model in R / a .stan file
 - Auto-rendering of C++/stan code

WORKFLOW USING STAN

1. Formulate the statistical model
 - Write your model in a `.stan` file
2. Sample from the posterior distribution
 - Run the `stan()` function in R
3. Assess (non)—convergence
 - Render plots and diagnostics
4. Draw inference

BASIC STAN PROGRAM

```
functions { //optional
}

data { //required
// define the format of your data
}

transformed data { //optional
// any manipulations to the data
}

parameters { //required
// define the parameters used in the model
}

transformed parameters { //optional
// define any transformations of the parameters
// define any transformed parameters you want to sample
}

model { //required
// specify the model (likelihood, priors) using (transformed) parameters
}

generated quantities{ //optional
// any secondary output, that does not rely on the model
}
```

EXAMPLE: STATISTICAL MODEL

Simple linear regression model
1 predictor (x), 1 outcome (y)

$$y = b_0 + b_1x + \epsilon$$

$$\hat{y} = b_0 + b_1x$$

$$\epsilon \sim N(0, \sigma^2)$$

$$y \sim N(\hat{y}, \sigma^2)$$

EXAMPLE: STAN FILE

```
data {  
  int N;  
  vector[N] y;  
  vector[N] x;  
}  
parameters { //required  
  real b0;  
  real b1;  
  real sigma;  
}  
transformed parameters { //optional  
  vector[N] theta;  
  theta = b0 + b1*x;  
}  
model { //required  
  y ~ normal(theta, sigma);  
}
```

NOTES ON STAN CODE

- Data: specify all the data you are feeding in to the model
- Parameters: specify all model parameters
- Transformed parameters: specify all functions/calculations of core parameters that you want sampled as well
- Model: the likelihood and any priors

```
data {  
  int<lower=0> N;  
  vector[N] y;  
  vector[N] x;  
}  
  
parameters { //required  
  real b0;  
  real b1;  
  real sigma;  
}  
  
transformed parameters { //optional  
  vector[N] theta;  
  theta = b0 + b1*x;  
}  
  
model { //required  
  y ~ normal(theta, sigma);  
}
```

NOTES ON STAN CODE

- Specify the type of any quantity in stan
 - int, real, vector, matrix; specify limits
- For loops can be used, indexing starts at 1
- End each statement with a semicolon
- Rstan 'checks' the correctness of your code
 - Not always easy to figure out where the mistake is.

```
data {  
  int<lower=0> N;  
  vector[N] y;  
  vector[N] x;  
}  
parameters { //required  
  real b0;  
  real b1;  
  real sigma;  
}  
transformed parameters { //optional  
  vector[N] theta;  
  theta = b0 + b1*x;  
}  
model { //required  
  for(i in 1:N){  
    y[i] ~ normal(theta[i], sigma);  
  }  
}
```

ALTERNATIVE FORMULATIONS

```
data {  
  int N;  
  vector[N] y;  
  vector[N] x;  
}  
parameters { //required  
  real b0;  
  real b1;  
  real sigma;  
}  
transformed parameters { //optional  
  vector[N] theta;  
  theta = b0 + b1*x;  
}  
model { //required  
  y ~ normal(theta, sigma);  
}
```

```
data {  
  
  vector[30] y;  
  vector[30] x;  
}  
parameters { //required  
  real b0;  
  real b1;  
  real<lower=0> sigma;  
}  
transformed parameters { //optional  
}  
model { //required  
  for(i in 1:30){  
    y[i] ~ normal(b0 + b1*x[i], sigma);  
  }  
}
```


ALTERNATIVE FORMULATIONS

```
data {  
  int N;  
  vector[N] y;  
  vector[N] x;  
}  
parameters { //required  
  real b0;  
  real b1;  
  real sigma;  
}  
transformed parameters { //optional  
  vector[N] theta;  
  theta = b0 + b1*x;  
}  
model { //required  
y ~ normal(theta, sigma);  
}
```

```
data {  
  vector[30] y;  
  vector[30] x;  
}  
parameters { //required  
  real b0;  
  real b1;  
  real<lower=0> sigma;  
}  
transformed parameters { //optional  
}  
model { //required  
vector[N] theta;  
theta = b0 + b1*x;  
sigma ~ cauchy(0,1);  
target += normal_lpdf(y| theta, sigma);  
}
```

NOTES ON STAN CODE

- Any quantity in ``parameters`` is sampled
- Any quantity in ``transformed parameters`` is sampled from the posterior
- Log posterior value is incremented for each model specification
- By omitting priors/likelihoods for parameters, by default improper priors are used

SAMPLING FROM THE POSTERIOR WITH STAN

```
1 library(rstan)
2
3 N <- 40
4 x <- rnorm(N, 100, 20)
5 res <- rnorm(N, 0, 10)
6 y <- 10 + x * 5 + res
7 exldata <- list("N" = N, "x" = x, "y" = y)
8
9 fit1 <- stan(
10   file = "example1.stan", # Stan program
11   data = exldata,        # named list of data
12   chains = 4,            # number of Markov chains
13   warmup = 1000,         # number of warmup iterations per chain
14   iter = 2000,           # total number of iterations per chain
15   cores = 1,             # number of cores (could use one per chain)
16   verbose = TRUE,        # show progress
17   refresh = 10,          # number of iterations to refresh for
18   pars = "theta",
19   include = FALSE
20 )
```

NOTES ON RUNNING RSTAN

- Chains
 - how many MCMC chains are you using?
 - used to assess mixing and convergence
 - 3-4 general amount
- Warmup
 - how many iterations should be considered as warmup?
 - these are to 'train' the sampler
 - not used in final samples for estimation
- Iterations
 - total number of iterations
 - Warmup : iteration ratio 2:1 / 1:1 / 1:2
 - at least 1000 iterations for final samples
- Cores
 - if you have a multicore computer, you can start each chain on an individual core
 - speeds up computing time
 - cannot split up a chain over multiple cores
- Refresh & verbose
 - display intermediate samples

```
1 library(rstan)
2
3 N <- 40
4 x <- rnorm(N, 100, 20)
5 res <- rnorm(N, 0, 10)
6 y <- 10 + x * 5 + res
7 exldata <- list("N" = N, "x" = x, "y" = y)
8
9 fit1 <- stan(
10   file = "example1.stan", # Stan program
11   data = exldata,         # named list of data
12   chains = 4,             # number of Markov chains
13   warmup = 1000,          # number of warmup iterations per chain
14   iter = 2000,            # total number of iterations per chain
15   cores = 1,              # number of cores (could use one per chain)
16   verbose = TRUE,         # show progress
17   refresh = 10,           # number of iterations to refresh for
18   pars = "theta",
19   include = FALSE
20 )
```

NOTES ON RUNNING RSTAN

- Seed
 - Set the seed that stan uses (set.seed in R does not work)
- Init
 - Provide initial values for the sampler
- Pars
 - Indicate for which parameters you want the sampled output

Advanced options available in control argument, with for example the adapt_delta and max_treedepth (NUTS/HMC settings)

```
1 library(rstan)
2
3 N <- 40
4 x <- rnorm(N, 100, 20)
5 res <- rnorm(N, 0, 10)
6 y <- 10 + x * 5 + res
7 exldata <- list("N" = N, "x" = x, "y" = y)
8
9 fit1 <- stan(
10   file = "example1.stan", # Stan program
11   data = exldata,         # named list of data
12   chains = 4,             # number of Markov chains
13   warmup = 1000,          # number of warmup iterations per chain
14   iter = 2000,            # total number of iterations per chain
15   cores = 1,              # number of cores (could use one per chain)
16   verbose = TRUE,         # show progress
17   refresh = 10,           # number of iterations to refresh for
18   pars = "theta",
19   include = FALSE
20 )
```

NOTES ON RUNNING RSTAN

- Using the `stan()` function uses the sampler
 - Alternative: `optimizing()`
- Directly using the stan model file here
 - Also option to first compile model code and then sample

```
1 library(rstan)
2
3 N <- 40
4 x <- rnorm(N, 100, 20)
5 res <- rnorm(N, 0, 10)
6 y <- 10 + x * 5 + res
7 exldata <- list("N" = N, "x" = x, "y" = y)
8
9 fit1 <- stan(
10   file = "example1.stan", # Stan program
11   data = exldata,        # named list of data
12   chains = 4,            # number of Markov chains
13   warmup = 1000,         # number of warmup iterations per chain
14   iter = 2000,           # total number of iterations per chain
15   cores = 1,             # number of cores (could use one per chain)
16   verbose = TRUE,        # show progress
17   refresh = 10,          # number of iterations to refresh for
18   pars = "theta",
19   include = FALSE
20 )
```

```
stanmod1 <- rstan::stan_model(file = "file.stan")
stanmod2 <- rstan::stan_model(model_code = "data{real y;} parameters{real mu;} model{mu~normal(y,1);}")

samps1 <- rstan::sampling(stanmod1, data = list(y = 0))
```

ASSESSING CONVERGENCE

- Rstan has a lot of inbuilt functions to assess the convergence of your posterior samples
- Sampled output gives rhat and effective sample size
- Plot additional output using:
 - `plot()`
 - `traceplot()`
 - `pairs()`
 - `stan_diag([[multiple options]])`
 - `stan_rhat()`
 - `stan_par(par = "par")`

ESTIMATION

```
> fit1
```

```
Inference for Stan model: example1.
```

```
4 chains, each with iter=2000; warmup=1000; thin=1;
```

```
post-warmup draws per chain=1000, total post-warmup draws=4000.
```

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
b0	-3.70	0.24	9.13	-21.58	-9.70	-3.66	2.22	14.95	1424	1
b1	5.15	0.00	0.09	4.97	5.09	5.15	5.21	5.34	1420	1
sigma	12.59	0.04	1.48	10.13	11.56	12.43	13.50	15.90	1594	1
lp__	-120.63	0.04	1.33	-124.07	-121.19	-120.28	-119.68	-119.15	1139	1

```
Samples were drawn using NUTS(diag_e) at Fri Sep 16 10:27:34 2022.
```

```
For each parameter, n_eff is a crude measure of effective sample size,  
and Rhat is the potential scale reduction factor on split chains (at  
convergence, Rhat=1).
```


TROUBLESHOOTING

- Divergent transitions?
 - Inspect the posterior diagnostic figures, assess whether maybe there is 1 or multiple problematic chains
 - Inspect the per chain samples
 - Adjust control parameter – adapt delta, maxtreedepth, objective tolerance...
- Bulk / tail ESS too low? / Rhat above 1 / No mixing of chains?
 - More iterations, using thinning
 - Changing seed / starting values may sometimes do the trick
 - Reparameterization of model (centering, transformations, etc) to reduce strong autocorrelations in sampler
- Inspect the posterior to see if results are making sense!!! All transformation done correctly?

TROUBLESHOOTING

- I made changes to the model but these are not reflected in the output
 - Did you correctly reload the model? Stan makes 'interim' model codes, so can sometimes need twice to get updated correctly

TROUBLESHOOTING

- Stan does not start sampling
 - There is a bug somewhere in your code!
 - Some type mismatch or incorrect use of a function
 - Go back to file, see if the rstan help helps you detect the possible source
 - Reevaluate the functions you are using: are the inputs of the correct type?
 - Use print statements and verbose = TRUE to get interim diagnostics.

```
TRANSLATING MODEL 'example1' FROM Stan CODE TO C++ CODE NOW.
```

```
SYNTAX ERROR, MESSAGE(S) FROM PARSER:
```

```
error in 'model75164303c80_example1' at line 12, column 7
```

```
-----  
10: }  
11: transformed parameters { //optional  
12:   vector theta;  
      ^  
13:   theta = b0 + b1*x;  
-----
```

```
PARSER EXPECTED: <vector length declaration: data-only integer expression in square brackets>
```

```
Error in stanc(file = file, model_code = model_code, model_name = model_name, :
```

```
  failed to parse Stan model 'example1' due to the above error.
```

```
11 ▾ transformed parameters { //optional  
12 ▾ vector theta;  
12 ▾   theta = b0 + b1*x;  
12 ▾   ^  
12 ▾   PARSER EXPECTED:  
15 ▾ model { //required  
16 ▾   y ~ normal(theta, sigma);  
17 ▾ }  
18
```

TROUBLESHOOTING

- Stan user guides
 - User guide, many examples and explanations
 - https://mc-stan.org/docs/2_18/stan-users-guide/index.html
 - Function references
 - Vector, unit, matrix based functions
 - Distributions
 - https://mc-stan.org/docs/2_20/functions-reference/index.html#overview
- Stack exchange // google
- Stan forums

EXAMPLE IN PRACTICE

- See live demonstration